

# Deep Reinforcement Learning for Sequential Decision Making: A Case Study in Energy Management



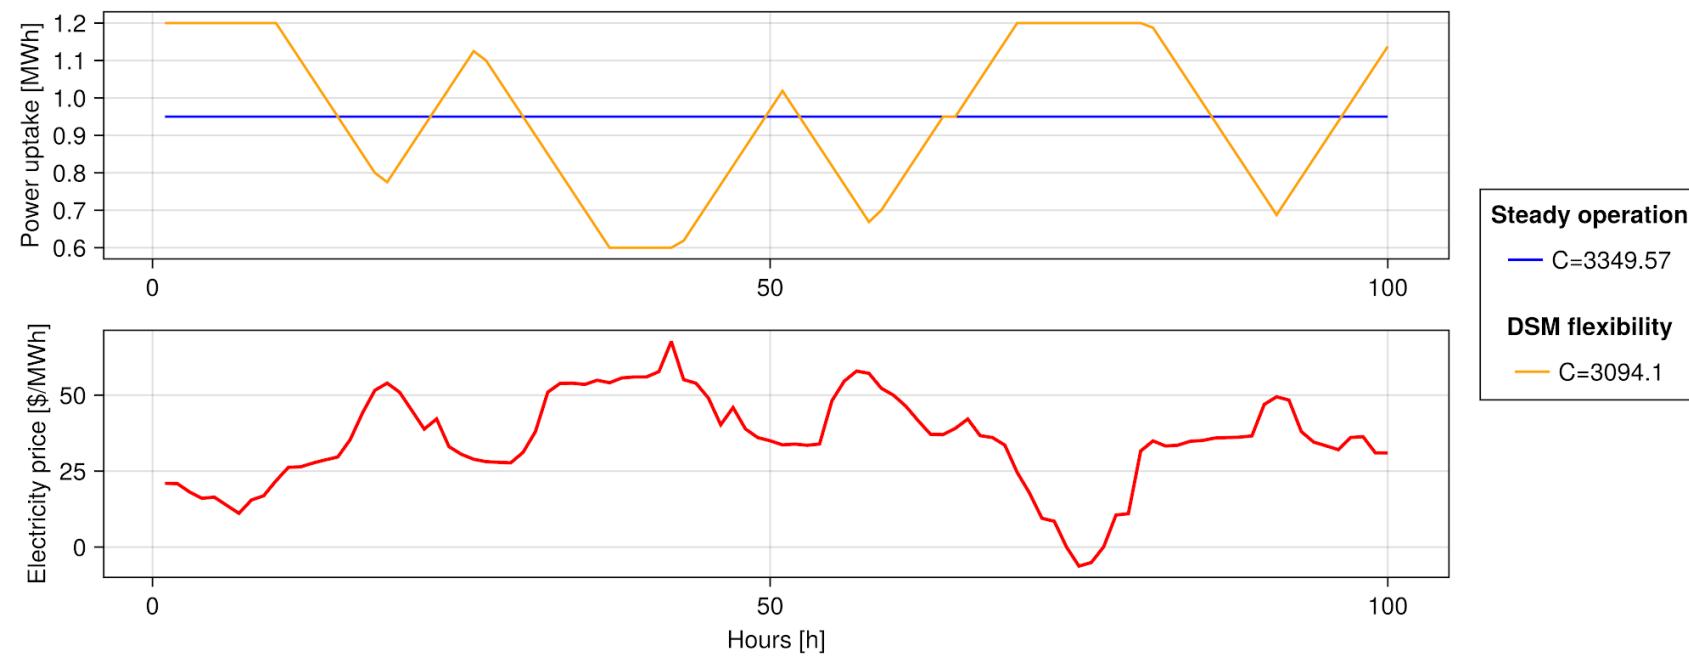
Zhe Li, Amin Nassaji, Alexandra Janey

April 22, 2025

# Problem Introduction

---

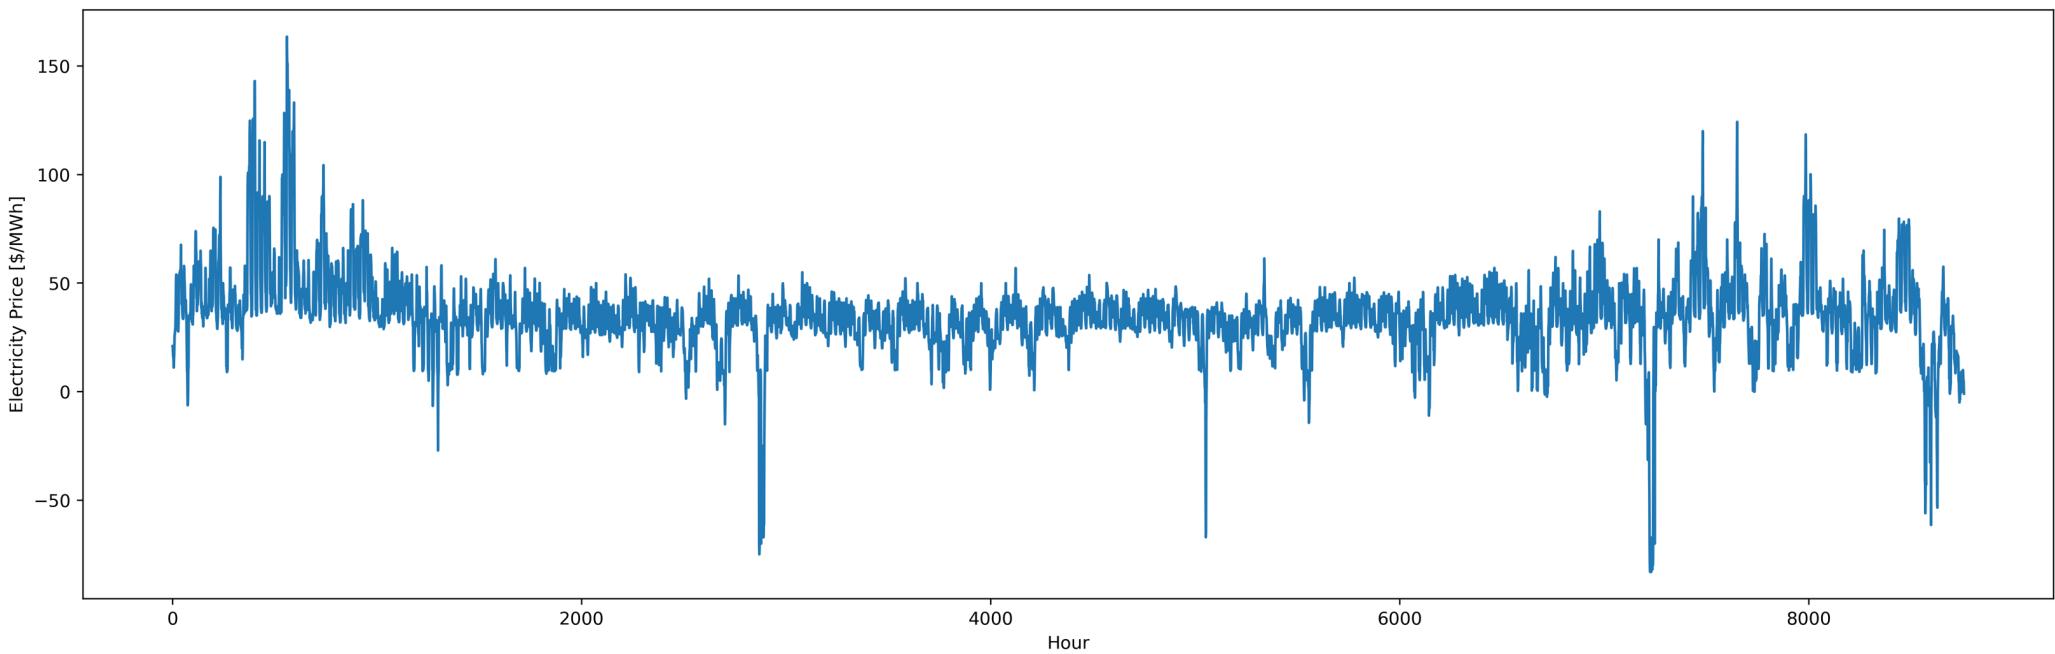
Electricity prices are volatile, so industrial demand side management (DSM) is a flexibility measure used for cost savings



# Dataset

---

German electricity sector, hourly spot prices in 2017

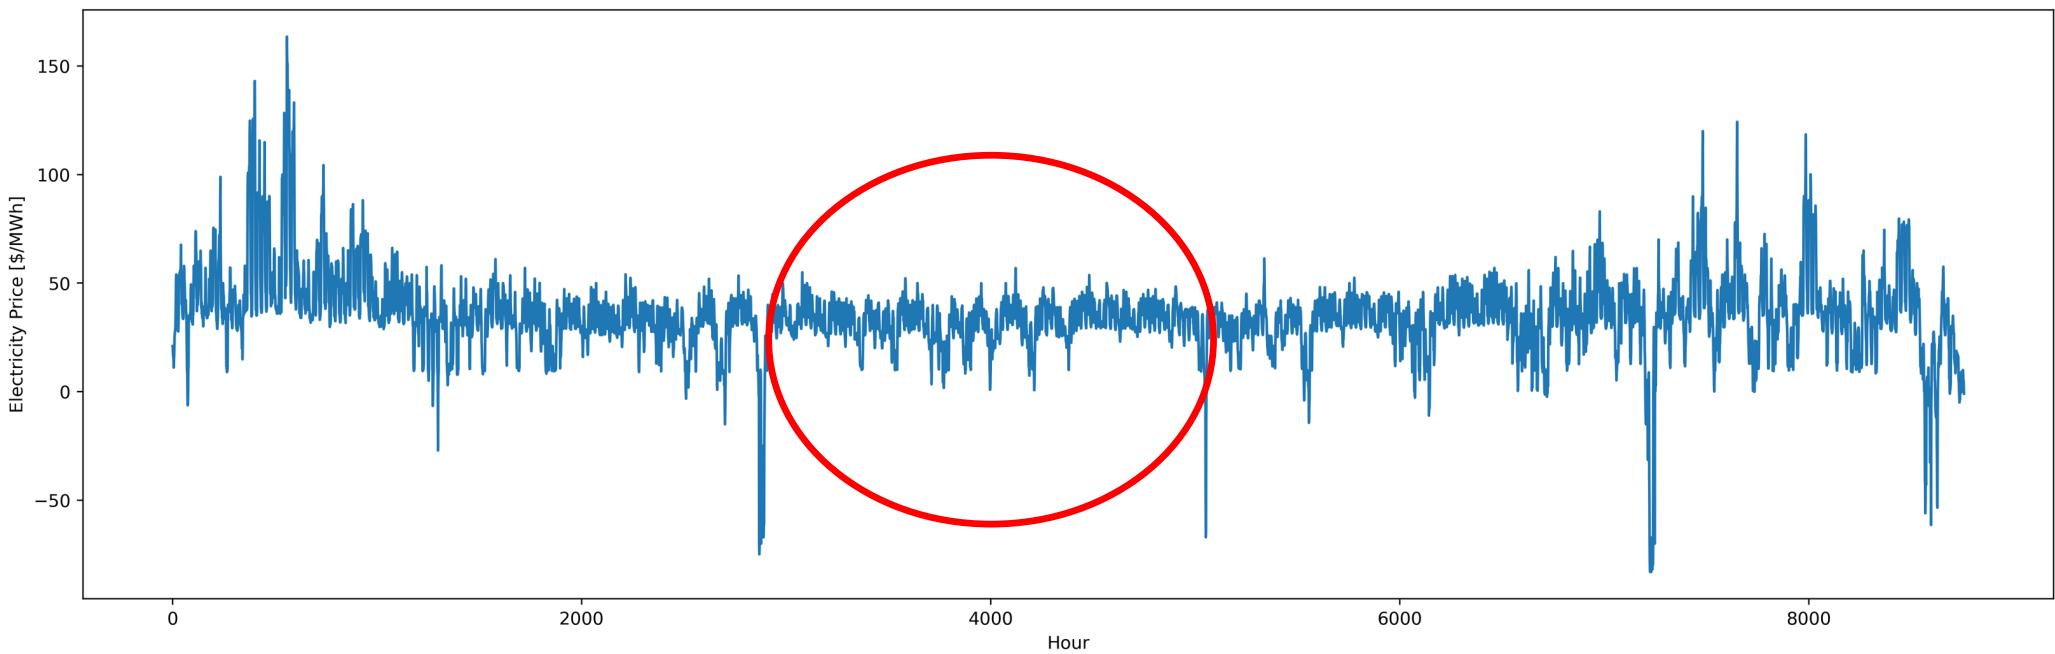


Source: <https://www.agora-energiewende.de/>

# Dataset

---

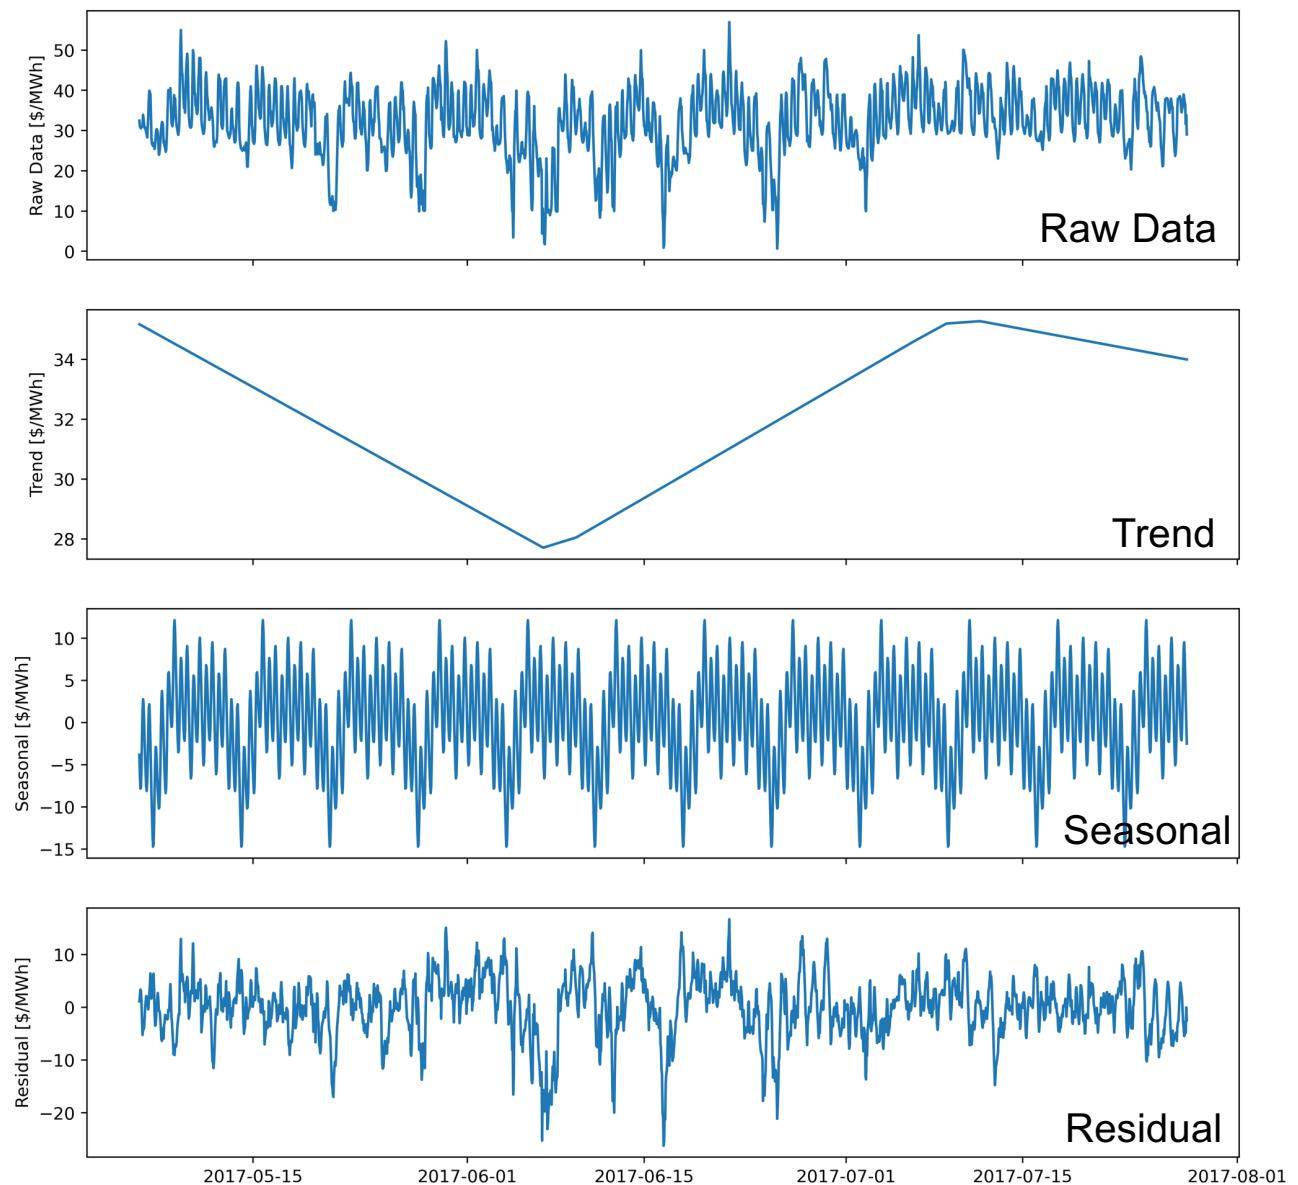
German electricity sector, hourly spot prices in 2017



Source: <https://www.agora-energiewende.de/>

# Time Series Analysis

Training data from  
hours 3000 – 5000  
(~May 6, 2017 – July 28, 2017)



Note: for test data we utilize  
hours 5880 - 6048

## Problem Description

---

- Traditional optimization methods have large computational costs when trying to implement in real time
- Treat scheduling model as a battery where we can “charge” it to store energy and use later
- Utilize sequential decision-making
- Solve this using Reinforcement Learning (RL) and compare it with the performance of Model Predictive Control (MPC)

## Deterministic Model (Oracle)

---

Main variables:  $c_t$  (cost),  $P_t$  (power uptake),  $S_t$  (storage level), and  $\Delta_t$  (ramping rate)

**minimize**  $C = \sum_{t \in \mathcal{N}} c_t \cdot P_t$  Minimize overall cost

**subject to**  $S_{t+1} = S_t + (P_t - P^{nom}) \quad \forall t \in \mathcal{N}$  Storage level

$$\Delta_{t+1} = P_{t+1} - P_t \quad \forall t \in \mathcal{N}$$
 Ramping rate

$$P^{nom} \cdot |\mathcal{N}| = \sum_{t \in \mathcal{N}} P_t$$
 Electricity demand

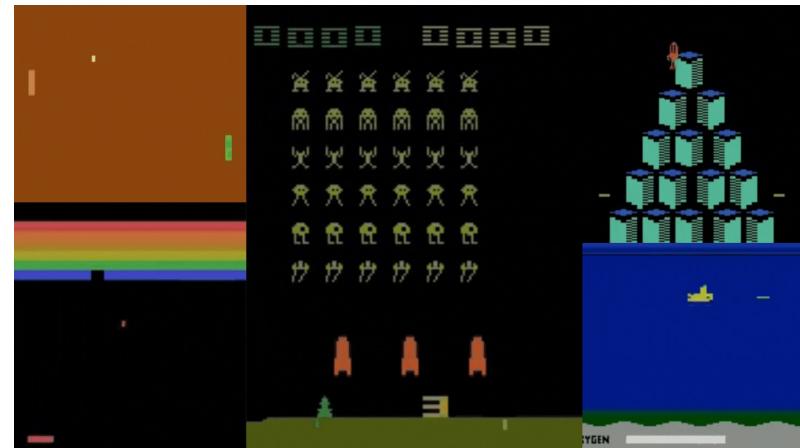
$$P^{\min} \leq P_t \leq P^{\max} \quad \forall t \in \mathcal{N}$$

$$-S^{\max} \leq S_t \leq S^{\max} \quad \forall t \in \mathcal{N}$$

$$-\Delta^{\max} \leq \Delta_t \leq \Delta^{\max} \quad \forall t \in \mathcal{N}$$

# Deep Reinforcement Learning: Overview

- What is RL?
  - Framework for solving sequential decision making problem
- How does it learn?
  - Trial and error from a world that provides occasional rewards
- What is the “Deep” part?
  - RL + neural network
- Applications
  - Games: Atari, Chess, Go...
  - Training large language model
  - Robotics controller

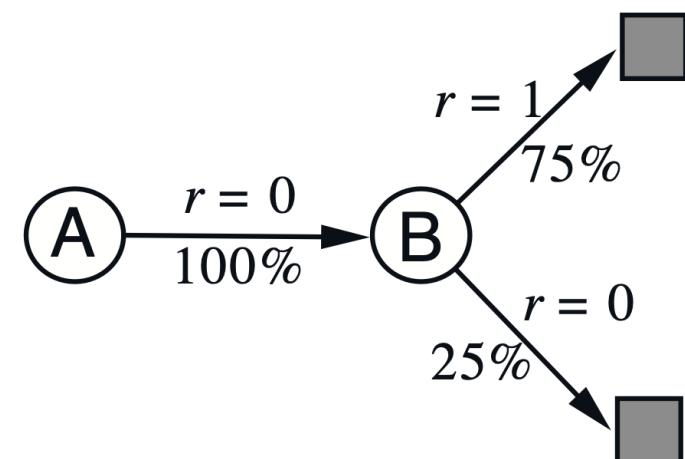


# Framework: Markov Decision Process (MDP)

- Sequential decision problem:

$$(s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_t, a_t, r_t, s_{t+1}, \dots)$$

- $s_t$ : **state** variable (what we know)
- $a_t$ : **action** variable (what we do)
- $c_t = c(s_t, a_t)$ : **cost** (what we get)
- $s' \sim p(s'|s, a)$ : **transition** dynamics
- $a_t \sim \pi(s_t)$ : **policy** (decision guidance)



# Markov Decision Process Modeling

---

- State variables (what we know)

$$s_t = (S_t, P_{t-1}, c_t, \Delta c_t)$$

- $S_t$ : Storage level at time t
- $P_{t-1}$ : Power uptake at time t-1
- $c_t$ : Electricity price at time t
- $\Delta c_t = c_t - c_{t-1}$ : Price change from time t-1 to t

- Action variables (what we do)

$$a_t = (P_t)$$

- $P_t \in [0.6, 1.2]$ : Power uptake at time t

# Markov Decision Process Modeling

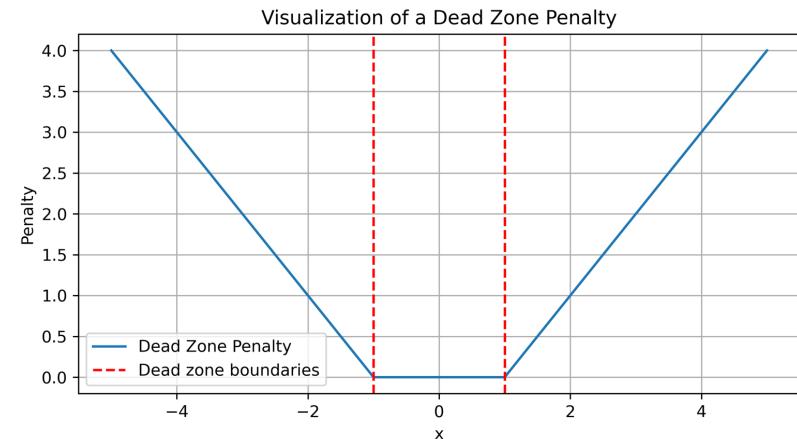
- Cost (what we get)

$$C_t = C(s_t, a_t) = c_t \cdot P_t + \text{Penalty}_{ramp} + \text{Penalty}_{storage}$$

- $c_t \cdot P_t$ : cost for purchasing electricity
- $\text{Penalty}_{ramp}(\Delta_t)$ : penalty for breaking ramping limit
- $\text{Penalty}_{storage}(S_t, P_t)$ : penalty for storage capacity

- Objective (cumulative cost):

$$\text{minimize } J = \sum_{t=0}^N C_t$$



# Markov Decision Process Modeling

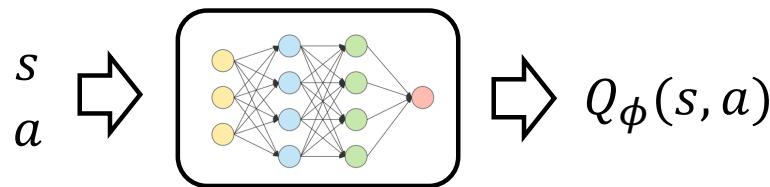
---

- Transition dynamics (how environment evolves)
  - Storage level:  $S_{t+1} \leftarrow S_t + (P_t - P^{nom})$
  - Spot electricity price:  $c_{t+1} \leftarrow \text{Historical Data}$
- Policy (what to do given the state)
  - Stochastic:  $a_t \sim \pi(a|s)$
  - Deterministic:  $a_t = \pi(s)$
- Task: Learn a good policy that minimizes cumulative cost  
**w/ Reinforcement Learning**

# Deep Reinforcement Learning

## Value-based RL

Learn a value function to evaluate action

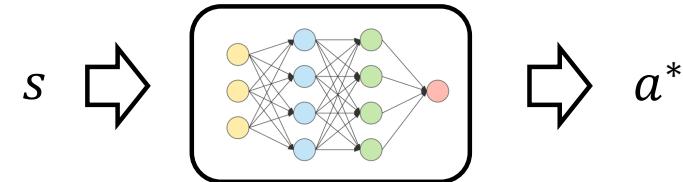


$$a^* = \operatorname{argmax}_a Q_\phi(s, a)$$

- Example: Q-learning
- Discrete action space

## Policy-based RL

Learn a direct mapping from state to action



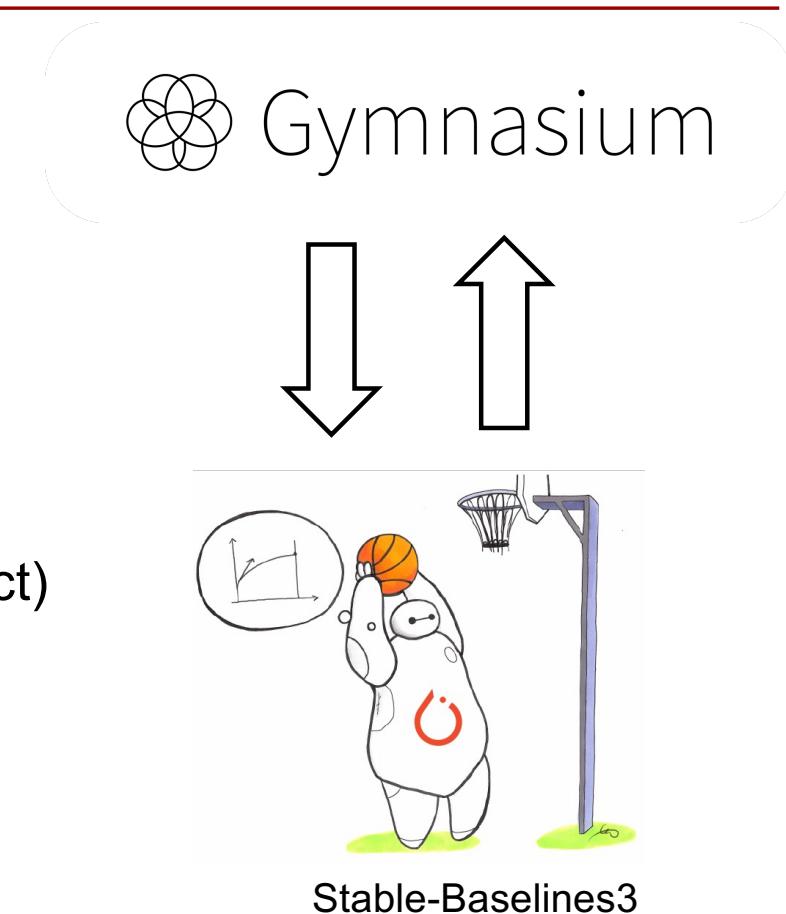
$$a^* = \pi_\theta(s)$$

- Example: Proximal Policy Opt.
- Continuous action space

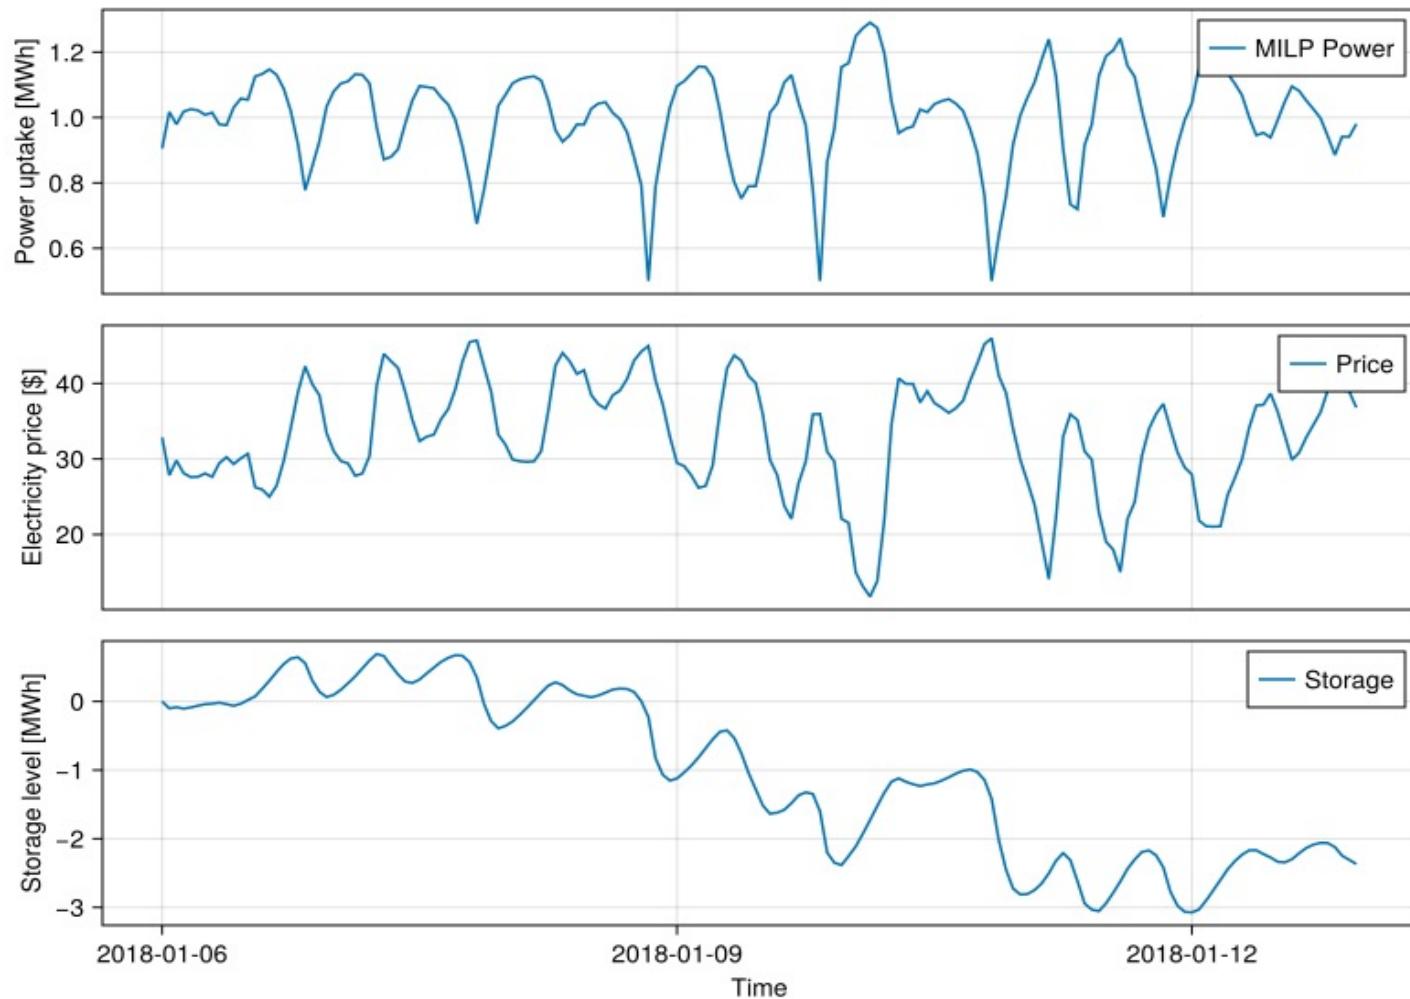
**Reminder:**  $s_t = (S_t, P_{t-1}, c_t, \Delta c_t)$        $a_t = (P_t)$

## Implementation Details

- Environment:
  - OpenAI's [Gymnasium](#) library
- Training Algorithm:
  - PPO from [Stable-Baselines3](#)
- Hyperparameters:
  - Policy network: 'MlpPolicy'
  - Gamma: 0.995 (weight on long-term effect)
  - Learning rate: 3e-4

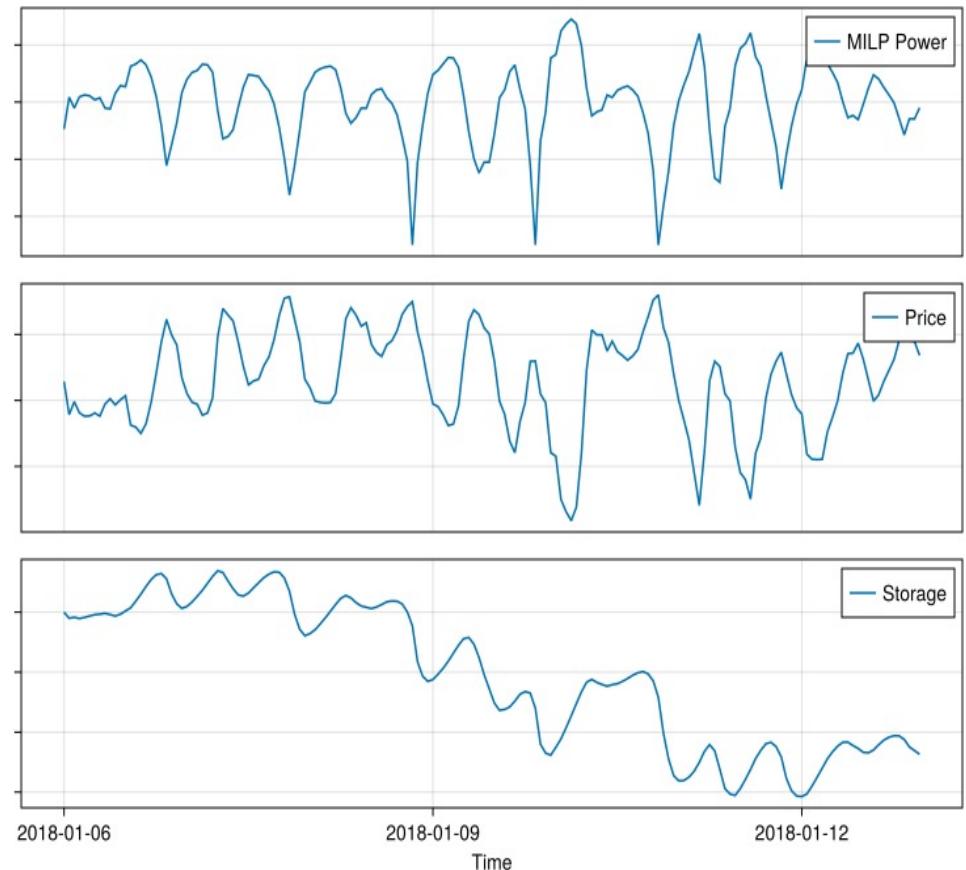


## Simulation results: MILP

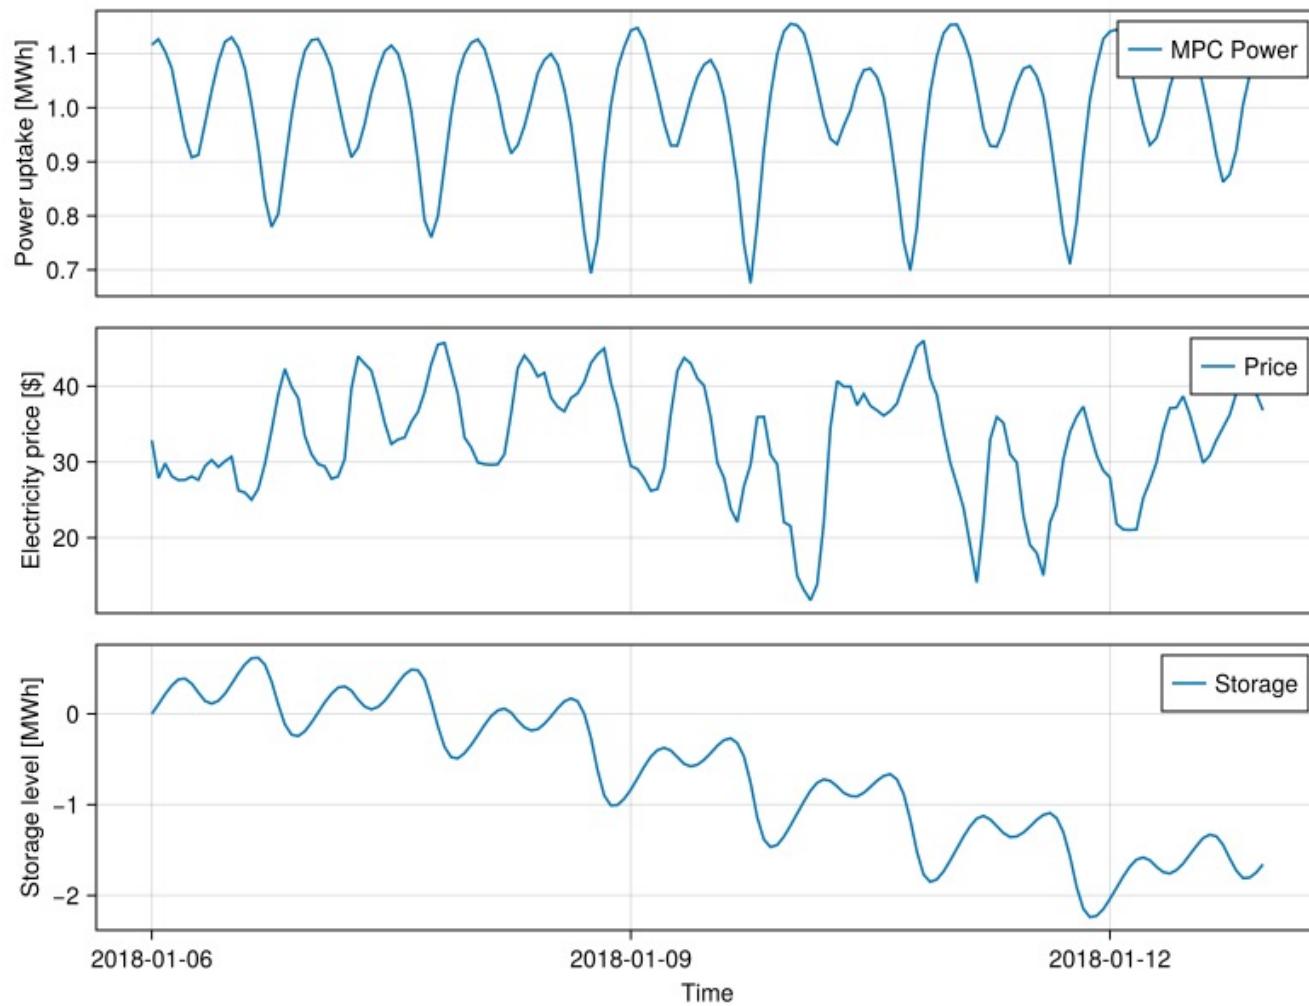


## MILP: Strengths and Trade-off

- Pros
  - Global Optimality
  - Precise Constraint Handling
  - Efficient for Planning
  
- Cons
  - Modeling Nonlinearity
  - Requires Forecasts
  - Binary Variables
  - Rigid Once Solved



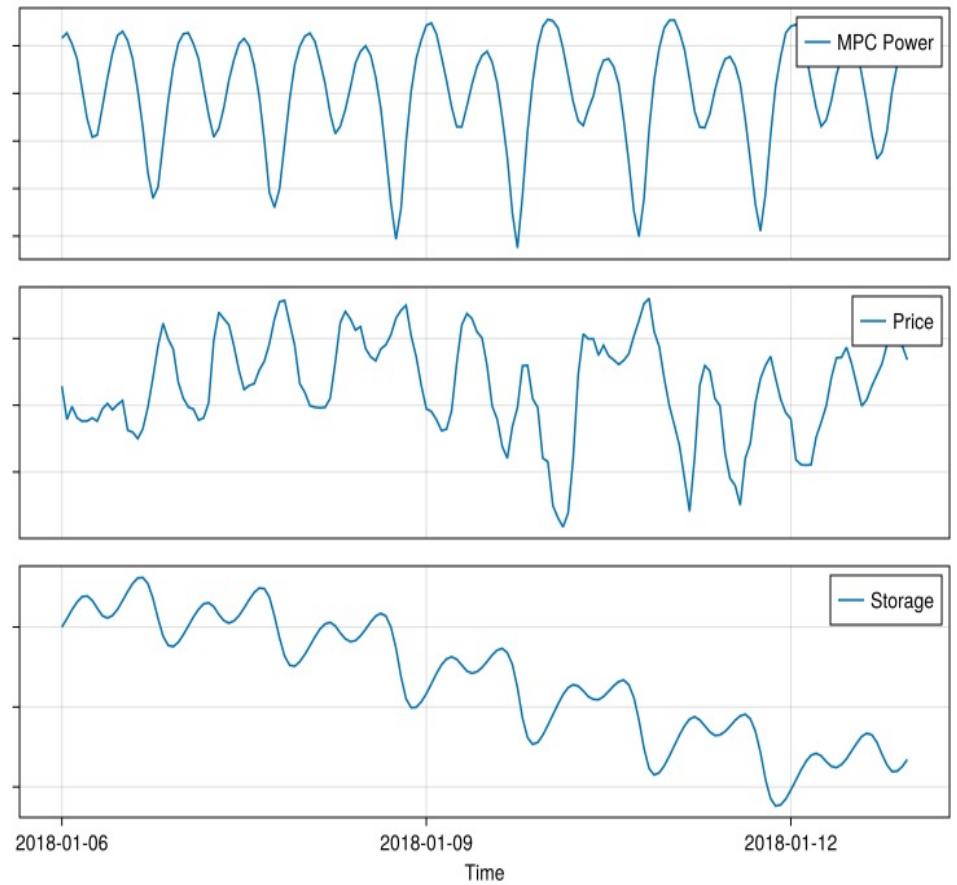
## Simulation results: MPC



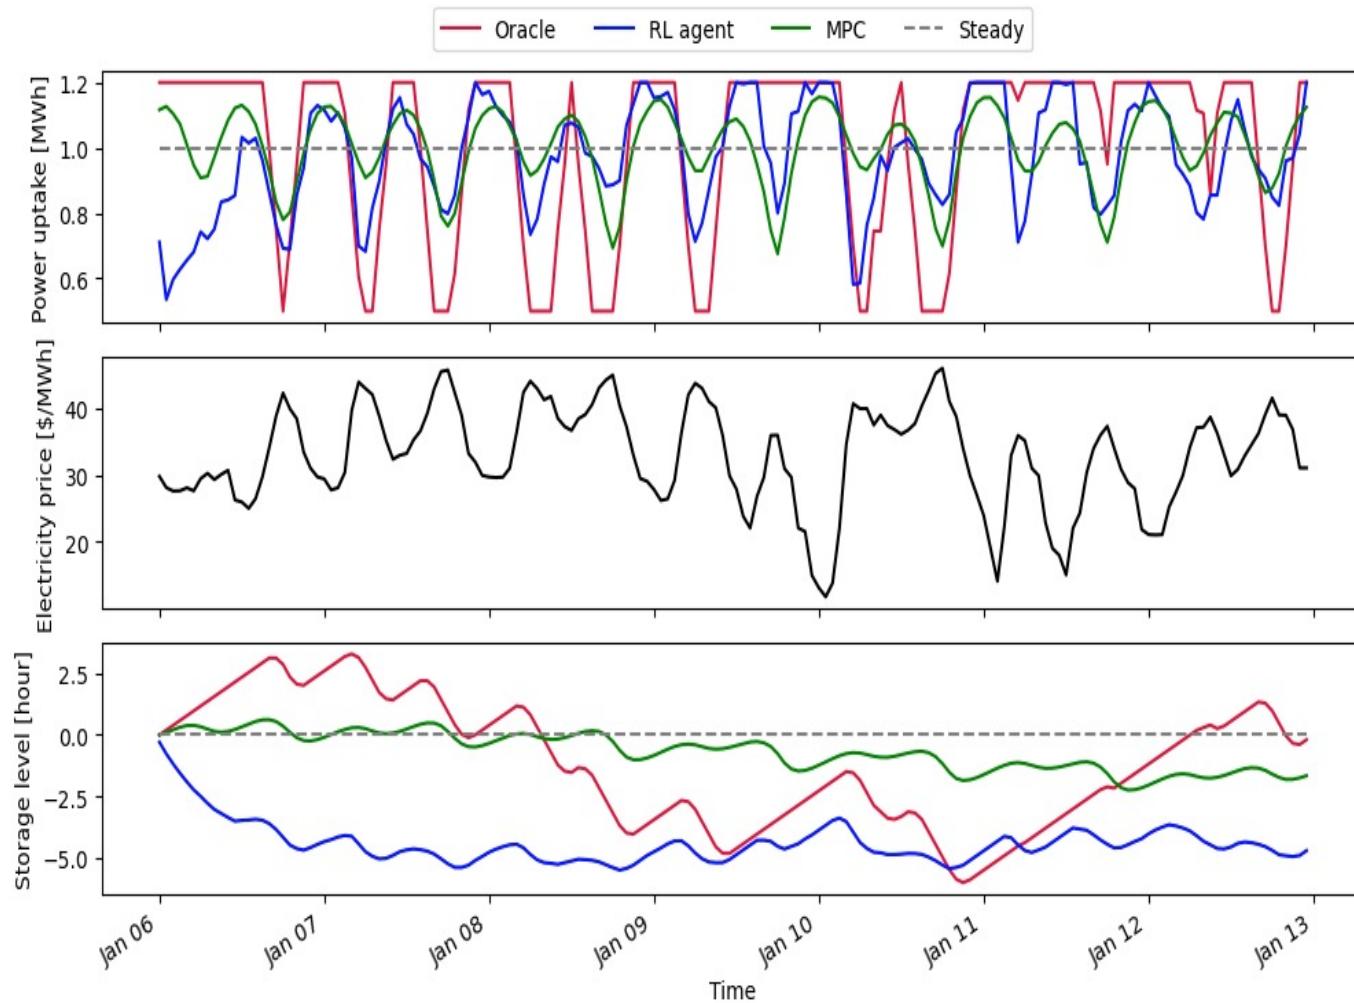
Total Cost:  
\$5462.73

## MPC: Strengths and Trade-off

- Pros
  - Real-Time Adaptivity
  - Handles Constraints Naturally
  - Smooth Operation
  - Robust to Forecast Horizon
  
- Cons
  - Forecast Dependency
  - Computational Load
  - No Global Optimality

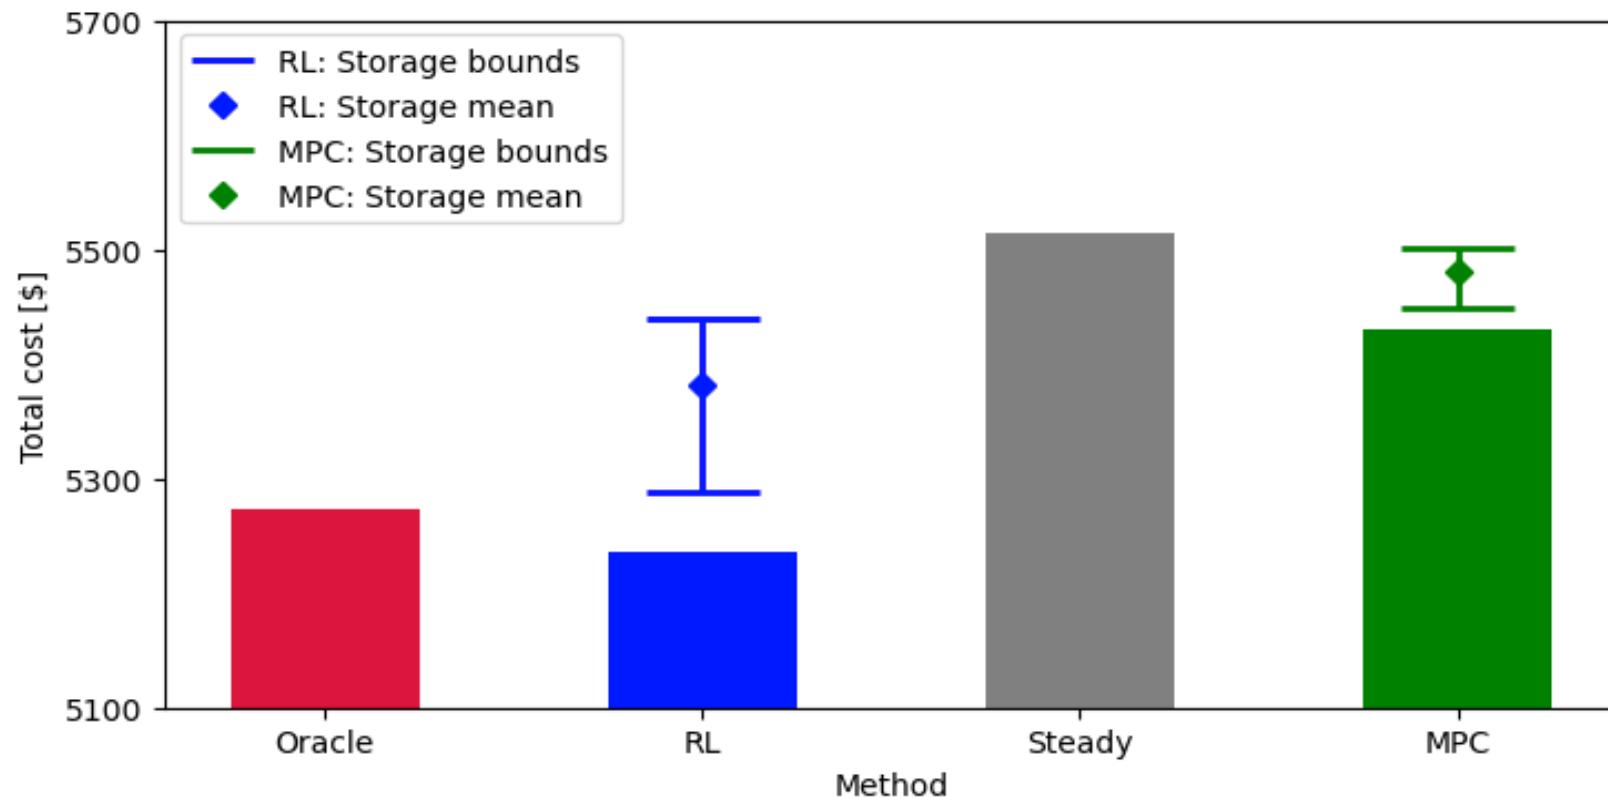


## Simulation results: RL



Total Cost:  
\$5236.00

## Final Comparison and Conclusion



## Final Comparison and Conclusion

---

- Pros
  - No Forecasting Required
  - Adaptivity
  - Smooth Behavior
  - Strong Cost Performance
  
- Cons
  - No Hard Guarantees
  - Training Time & Tuning
  - Less Transparent

